

COM6521 21/22 Assignment

ACP21MFM

1 OpenMP

1.1 Stage 1

In stage 1 (Tile Sum), different from the provided algorithm, the double tile loop is fixed into a single loop through index manipulation withing the loop to improve performance and make the code more readable. the original most inner loop is also replaced with assigning values directly to the mosaic sum since the original loop only served changing ch value, which was always 3. OpenMP scoping was kept simple due to the reason that the variables used inside the parallel region were defined in the region, which makes them private to each thread regardless if they were explicitly scoped or not. through experimentation, this had a slight improvement on the performance.

```
#pragma omp parallel for private(t, p_x, p_y)
for (t = 0; t < omp_TILES_X * omp_TILES_Y; ++t) {
    // get original loop indecies to traverse the image
    const int t_x = t % omp_TILES_X;
    const int t_y = t / omp_TILES_X;

    const unsigned int tile_index = t * omp_input_image.channels;
    const unsigned int tile_offset = (t_y * omp_TILES_X * TILE_SIZE * TILE_SIZE + t_x * TILE_SIZE) * omp_input_image.channels;

    // For each pixel within the tile
    for (p_x = 0; p_x < TILE_SIZE; ++p_x) {
        for (p_y = 0; p_y < TILE_SIZE; ++p_y) {
            // For each colour channel
            const unsigned int pixel_offset = (p_y * omp_input_image.width + p_x) * omp_input_image.channels;

            // Avoiding loop usage since channels are always 3 (R, G, B), improves performance and code coherency
            omp_mosaic_sum[tile_index] += omp_input_image.data[tile_offset + pixel_offset];
            omp_mosaic_sum[tile_index + 1] += omp_input_image.data[tile_offset + pixel_offset + 1];
            omp_mosaic_sum[tile_index + 2] += omp_input_image.data[tile_offset + pixel_offset + 2];
        }
    }
}
```

Figure 1: OpenMP Stage 1 Code.

1.2 Stage 2

Stage 2, The average pixel value for each tile is calculated by dividing the tile sums from the previous stage by the square of TILE SIZE (the number of pixels in a tile). Such a task is considered very simple even for a large buffer of pixel values and thus the task of paralleling this stage was perplexing. Many methods were tried in order to improve this stage's time but the original sequential algorithm gave faster times.

The provided solution is given as an example of what might make this stage faster for a harder problem, in order to use OpenMP's reduction method in the summation of the pixels, three variables were defined, each representing a different channel (R, G, B). calculated in the parallel section similarly to the original algorithm, then combined and divided after the parallel section in output global average, with respect to each channel's index. While the solution still provided slower performance, through experimentation, it was the fastest.

The initial speculation that this is happening is because the simple calculation requirements, the parallel approach is disadvantageous to the sequential approach for memory allocation reasons.

In real life scenarios, such a piece of code won't be integrated into parallel code since the sequential approach is faster, which is the whole purpose of parallel programming.

```

int t;

int sum_r = 0;
int sum_g = 0;
int sum_b = 0;

#pragma omp parallel for private(t) reduction(+: sum_r, sum_g, sum_b)
for (t = 0; t < totalTiles; ++t) {
    const unsigned int tile_index = t * omp_input_image.channels;

    omp_mosaic_value[tile_index] = (unsigned char)(omp_mosaic_sum[tile_index] / TILE_PIXELS); // Integer division is fine here
    omp_mosaic_value[tile_index + 1] = (unsigned char)(omp_mosaic_sum[tile_index + 1] / TILE_PIXELS);
    omp_mosaic_value[tile_index + 2] = (unsigned char)(omp_mosaic_sum[tile_index + 2] / TILE_PIXELS);

    sum_r += omp_mosaic_value[tile_index];
    sum_g += omp_mosaic_value[tile_index + 1];
    sum_b += omp_mosaic_value[tile_index + 2];
}

// divide by tile size to get average and recombine to output_global_average
output_global_average[0] = (unsigned char)(sum_r / (totalTiles));
output_global_average[1] = (unsigned char)(sum_g / (totalTiles));
output_global_average[2] = (unsigned char)(sum_b / (totalTiles));

```

Figure 2: OpenMP Stage 2 Code.

1.3 Stage 3

In stage 3, the final output image is derived by broadcasting the pixel averages from the compact mosaic to a full scale image.

Similar to the approach in Stage 1, the tile loops in this stage are fixed into a single loop through index manipulation, which improves overall performance and makes OpenMP's operations and scoping less complicated.

One more difference from the original algorithm, is removing the usage of memcpy function to assign the pixel values to the output. in C/C++, memcpy is considered to be IO- operation bound is not bound to the CPU so applying parallel programming techniques to it is not applicable or useful, instead, assigning pixel values directly to each distinct channel (indexes [N, N+1, N+2]) is more parallel- friendly and enhanced performance considerably since such an operation is a processing-unit-bound, thus, multi thread operations are able to benefit it.

```

int t, p_x, p_y;

#pragma omp parallel for private(t, p_x, p_y)
// For each tile
for (t = 0; t < omp_TILES_X * omp_TILES_Y; ++t) {
    // get original loop indecies to map values to output
    const int t_x = t % omp_TILES_X;
    const int t_y = t / omp_TILES_X;

    const unsigned int tile_index = t * omp_input_image.channels;
    const unsigned int tile_offset = (t_y * omp_TILES_X * TILE_SIZE * TILE_SIZE + t_x * TILE_SIZE) * omp_input_image.channels;

    // For each pixel within the tile
    for (p_x = 0; p_x < TILE_SIZE; ++p_x) {
        for (p_y = 0; p_y < TILE_SIZE; ++p_y) {
            const unsigned int pixel_offset = (p_y * omp_input_image.width + p_x) * omp_input_image.channels;
            // Copy whole pixel
            omp_output_image.data[tile_offset + pixel_offset] = omp_mosaic_value[tile_index];
            omp_output_image.data[tile_offset + pixel_offset + 1] = omp_mosaic_value[tile_index + 1];
            omp_output_image.data[tile_offset + pixel_offset + 2] = omp_mosaic_value[tile_index + 2];
        }
    }
}

```

Figure 3: OpenMP Stage 3 Code.

2 CUDA

2.1 Stage 1 Host Code

For CUDA paralleling, the host specifies the kernel parameters, and usually allocates memory for the passed variables. The main required kernel parameters are 1) how many blocks per grid to utilize, blocks_per_grid 2) how many threads per block to utilize, threads_per_block. both parameters' values are chosen, usually, to maximize processor occupancy, thus, increasing performance. in our case block number is chosen according to how many tiles we have in both directions, and thread number is chosen according to tile_size, which makes each pass of an image array occupy a full grid.

After execution of the Kernel is finished, usually, the calculated variable is retrieved to the host via cudaMemcpy, however, since d_mosaic_sum will be used on the device in later stage cudaMemcpy was skipped in this stage, which contributed to performance.

```
dim3 blocks_per_grid(cuda_TILES_X, cuda_TILES_Y, 1);
dim3 threads_per_block(TILE_SIZE, TILE_SIZE, 1);

kernel_stage1 <<<blocks_per_grid, threads_per_block >>>(d_input_image_data, d_mosaic_sum);

// No Need to retrieve mosaic_sum since we can use d_mosaic_sum
```

Figure 4: Stage 1 Host Code.

2.2 Stage 1 Kernel

In consistency with block and threads size chosen by the host, the original algorithm's loop parameters are chosen as each tile starting with the block ID and each pixel chosen according to thread ID. and just like the approach in OpenMP, array usage is avoided through breaking down into three different variables (R, G, B) which eliminates channel offset usage as well. Finally, in order to avoid branch divergence conditions __syncthreads() was used which improved performance by 30% for stage 1.

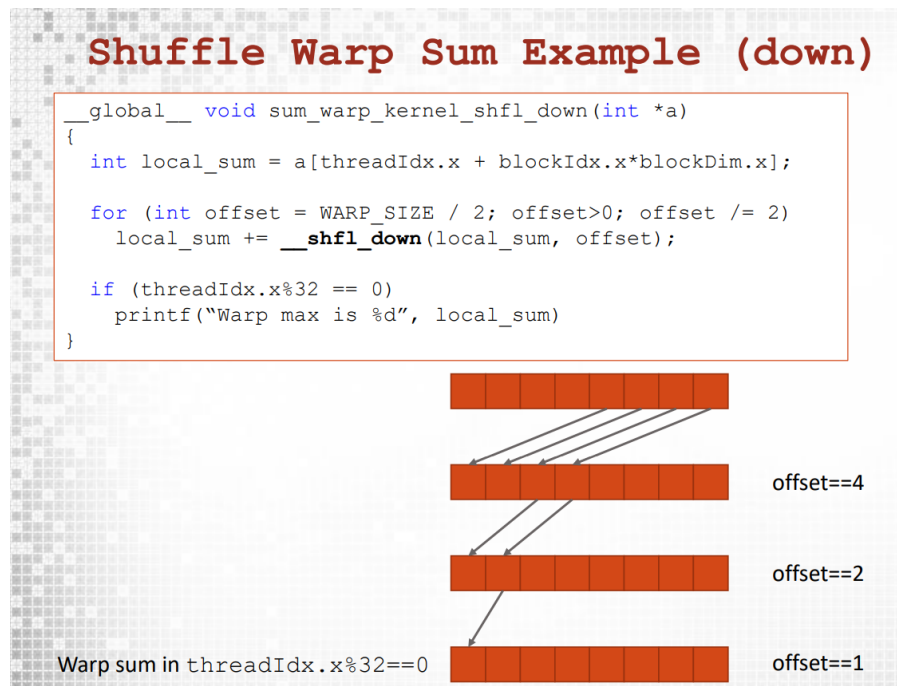


Figure 5: __shfl_down() logic, From Lecture12 COM6521 By Dr Paul Richmond.

For the summation of the pixel values, wrap operation shuffling down was used, which was the best one for the provided algorithm as its application purpose directly goes with the task as hand, moreover, to recombine the values into the mosaic sum array, atomicAdd was used to avoid race conditions.

```

__global__ void kernel_stage1(unsigned char* d_input_image_data, unsigned long long* d_mosaic_sum) {

    unsigned int t_x = blockIdx.x;
    unsigned int t_y = blockIdx.y;
    unsigned int p_x = threadIdx.x;
    unsigned int p_y = threadIdx.y;

    const unsigned int tile_index = (t_y * d_TILES_X + t_x) * d_input_image_channels;
    const unsigned int tile_offset = (t_y * d_TILES_X * TILE_SIZE * TILE_SIZE + t_x * TILE_SIZE) * d_input_image_channels;
    const unsigned int pixel_offset = (p_y * d_input_image_width + p_x) * d_input_image_channels;

    unsigned int r_sum = d_input_image_data[tile_offset + pixel_offset];
    unsigned int g_sum = d_input_image_data[tile_offset + pixel_offset + 1];
    unsigned int b_sum = d_input_image_data[tile_offset + pixel_offset + 2];

    for (int offset = 16; offset > 0; offset /= 2) {
        r_sum += __shfl_down(r_sum, offset);
        g_sum += __shfl_down(g_sum, offset);
        b_sum += __shfl_down(b_sum, offset);
    }

    if (threadIdx.x % 32 == 0) {
        //avoiding loop usage since channels are always 3 (R, G, B), improves performance and code coherency
        atomicAdd(&d_mosaic_sum[tile_index], r_sum);
        atomicAdd(&d_mosaic_sum[tile_index + 1], g_sum);
        atomicAdd(&d_mosaic_sum[tile_index + 2], b_sum);
    }

    //Avoid branch divergence from above
    __syncthreads();
}

```

Figure 6: Stage 1 Kernel.

2.3 Stage 2 Host Code

Just like stage 1, in order to maximize hardware occupancy, thread and block size were calculated. in most architectures the maximum thread per block number is 1024, which in our case directly goes with the tile size (ends up with grid size of 1 and thread number of 1024), in such a way, if a bigger tile size is provided the parameters would utilize the kernel in the most efficient way too. However, like OpenMP, CUDA failed to deliver better performance as the task is relatively simple.

The provided solution was found to be the fastest through experimentation when forcing CUDA implementation. As mentioned before this contradicts the purpose of parallel programming and in real life scenarios parallelism would not be applied here.

Afterwards, the global average is retrieved to the host through cudaMemcpy and the global sums are recombined and divided by total tiles to gain the average.

```

unsigned int total_tiles = cuda_TILES_X * cuda_TILES_Y;
unsigned int grid_size = (unsigned int)ceil(((double)(total_tiles)) / 1024);
unsigned int threads_n = (unsigned int)ceil((double)total_tiles / grid_size);

dim3 blocks_per_grid(grid_size);
dim3 threads_per_block(threads_n);

kernel_stage2<<<blocks_per_grid, threads_per_block >>>(d_mosaic_sum, d_mosaic_value, d_global_pixel_sum);

CUDA_CALL(cudaMemcpy(cuda_global_pixel_sum, d_global_pixel_sum, cuda_input_image.channels * sizeof(unsigned long long), cudaMemcpyDeviceToHost));

//Recombine into main host variable
output_global_average[0] = (unsigned char)(cuda_global_pixel_sum[0] / (total_tiles));
output_global_average[1] = (unsigned char)(cuda_global_pixel_sum[1] / (total_tiles));
output_global_average[2] = (unsigned char)(cuda_global_pixel_sum[2] / (total_tiles));

```

Figure 7: Stage 2 Host Code.

2.4 Stage 2 Kernel

Similarly to stage 1, loop usage is avoided through directly indexing the specific channel, (in our application this would always work, but it's assumed that RGB images are used and the indexing is always set up in the same approach).

```

__global__ void kernel_stage2(unsigned long long* d_mosaic_sum, unsigned char* d_mosaic_value, unsigned long long* d_global_pixel_sum) {
    unsigned int t = blockIdx.x * blockDim.x + threadIdx.x;

    d_mosaic_value[t * d_input_image_channels] = (unsigned char)(d_mosaic_sum[t * d_input_image_channels] / TILE_PIXELS);
    d_mosaic_value[t * d_input_image_channels + 1] = (unsigned char)(d_mosaic_sum[t * d_input_image_channels + 1] / TILE_PIXELS);
    d_mosaic_value[t * d_input_image_channels + 2] = (unsigned char)(d_mosaic_sum[t * d_input_image_channels + 2] / TILE_PIXELS);

    unsigned int r_sum = d_mosaic_value[t * d_input_image_channels];
    unsigned int g_sum = d_mosaic_value[t * d_input_image_channels + 1];
    unsigned int b_sum = d_mosaic_value[t * d_input_image_channels + 2];

    for (int offset = 16; offset > 0; offset /= 2) {
        r_sum += __shfl_down(r_sum, offset);
        g_sum += __shfl_down(g_sum, offset);
        b_sum += __shfl_down(b_sum, offset);
    }

    if (threadIdx.x % 32 == 0) {
        atomicAdd(&d_global_pixel_sum[0], r_sum);
        atomicAdd(&d_global_pixel_sum[1], g_sum);
        atomicAdd(&d_global_pixel_sum[2], b_sum);
    }
}

```

Figure 8: Stage 2 Kernel.

2.5 Stage 3 Host Code

Stage 3 host code can be approached like stage 1, as the same number of threads and blocks is the most suitable for the task.

```

dim3 blocks_per_grid(cuda_TILES_X, cuda_TILES_Y, 1);
dim3 threads_per_block(TILE_SIZE, TILE_SIZE, 1);

```

Figure 9: Stage 3 Host Code.

2.6 Stage 3 Kernel

The kernel code is relatively simple, thread and block locations' IDs are utilized to eliminate loop usage and broadcasting is done with the provided tile and pixel offset variables.

```

__global__ void kernel_stage3(unsigned char* d_output_image_data, unsigned char* d_mosaic_value) {
    unsigned int t_x = blockIdx.x;
    unsigned int t_y = blockIdx.y;
    unsigned int p_x = threadIdx.x;
    unsigned int p_y = threadIdx.y;

    const unsigned int tile_index = (t_y * d_TILES_X + t_x) * d_input_image_channels;
    const unsigned int tile_offset = (t_y * d_TILES_X * TILE_SIZE * TILE_SIZE + t_x * TILE_SIZE) * d_input_image_channels;
    const unsigned int pixel_offset = (p_y * d_input_image_width + p_x) * d_input_image_channels;

    d_output_image_data[tile_offset + pixel_offset] = d_mosaic_value[tile_index];
    d_output_image_data[tile_offset + pixel_offset + 1] = d_mosaic_value[tile_index + 1];
    d_output_image_data[tile_offset + pixel_offset + 2] = d_mosaic_value[tile_index + 2];
}

```

Figure 10: Stage 3 Kernel.

3 Results

From the results below, it is apparent that parallel programming techniques are greatly beneficial to increasing speeds of such algorithms, in some cases CUDA, for example, improved performance by more than 100x. Also, it was noted that parallel methods work best with computationally expensive operations, or cpu-bound operations. Using all of the techniques that were learned during the course of COM6521, parallel methods were successfully applied in the assignment, while also, discovering the best case uses for parallel programming.

if time was not a constraint, more can be discovered using profiling and benchmarking methods (Nsight compute, Visual Profiler) to improve the speeds of CUDA and aslo better algorithms could be developed to use within OpenMP boundaries as the provided algorithm was developed thinking in sequential manner.

256x256	CPU	OpenMP	CUDA
Stage 1	0.148ms	0.100ms	0.021ms
Stage 2	0.007ms	0.028ms	0.056ms
Stage 3	0.193ms	0.029ms	0.015ms

Figure 11: Results on 256x256 picture.

1024x1024	CPU	OpenMP	CUDA
Stage 1	2.241ms	0.367ms	0.105ms
Stage 2	0.011ms	0.050ms	0.059ms
Stage 3	3.029ms	0.512ms	0.057ms

Figure 12: Results on 1024x1024 picture.

2048x2048	CPU	OpenMP	CUDA
Stage 1	10.212ms	1.238ms	0.393ms
Stage 2	0.021ms	0.036ms	0.052ms
Stage 3	19.458ms	1.870ms	0.203ms

Figure 13: Results on 2048x2048 picture.

4096x4096	CPU	OpenMP	CUDA
Stage 1	43.833ms	4.591ms	1.491ms
Stage 2	0.071ms	0.069ms	0.059ms
Stage 3	101.8ms	9.402ms	0.775ms

Figure 14: Results on 4096x4096 picture.

4 Nsight Compute

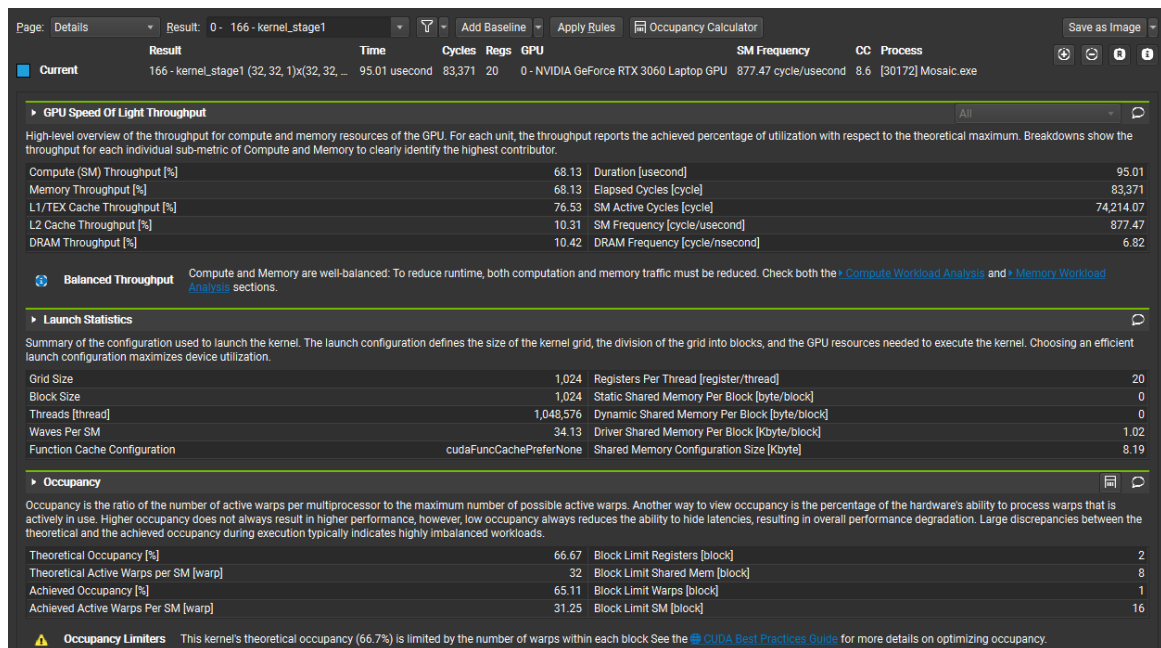


Figure 15: Stage 1 Nsight Analysis.

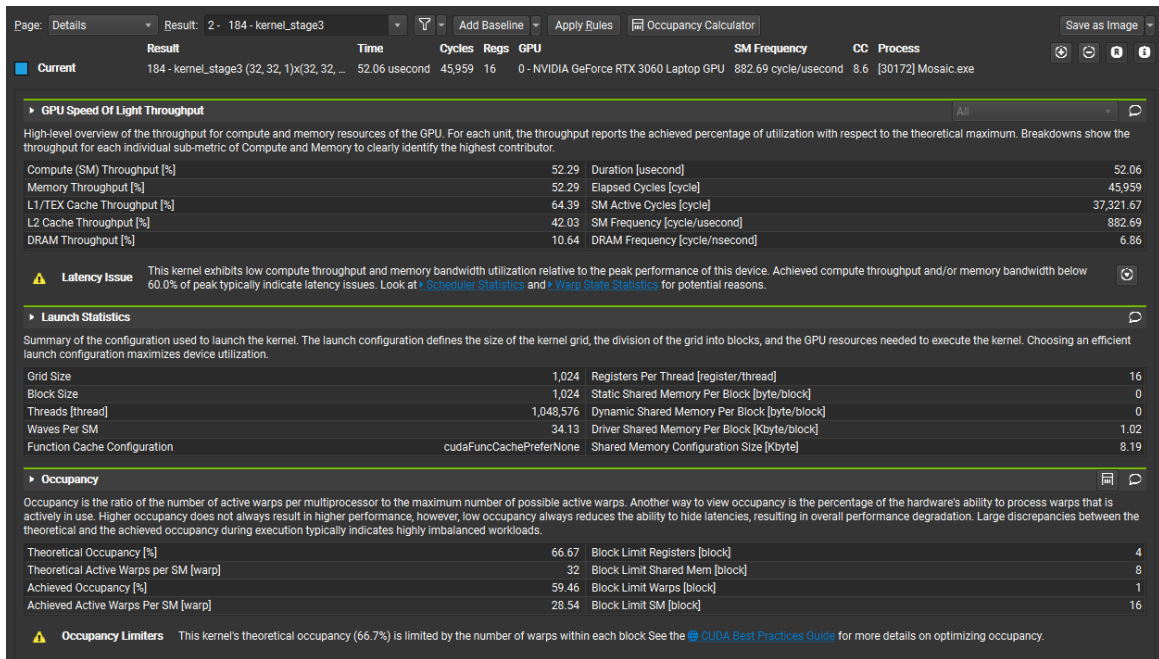


Figure 16: Stage 3 Nsight Analysis.